OCOM5300M Artificial Intelligence Project Conceptual Neuro-Symbolic Conformal Pipeline

In [1]:

```
# custom classes
from conceptual_neurosymbolic_conformal_pipeline import (
    ConceptIndex,
    ROCOv2_ConceptDataset,
    ConceptClassifierModel,
    UMLSCUISymbolicRuleGenerator,
    UMLSSymbolicValidator,
    ConformalCalibrator,
    ConceptualNeuroSymbolicConformalPipeline
)

# deep learning Functionality
import torch  # PyTorch library

# batching, shuffling and data loading capabilities
from torch.utils.data import DataLoader

# data handling and analysis
import pandas as pd

# utilities
import os  # file and directory operations

# file and directory paths
from pathlib import Path

# handling csv files
import csv

# frequencies
from collections import Counter

# pretrained Medical Model, i.e., OpenCLIP implementation used to load BiomedCLIP
import open_clip
```

In [2]:
```
# suppress Intel specific Math Kernel Library (MKL) warnings
os.environ['KMP_DUPLICATE_LIB_OK'] = 'True'
```

In [3]:
```
# detect device (CPU/GPU)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

In [4]:
```
# display the complete content of DataFrame columns without truncation
pd.set_option('display.max_colwidth', None);
```

In [5]:
```
# set random seed to 42
seed = 42
```

In [6]:
```
# set PyTorch random number generator seed for reproducible results
torch.manual_seed(seed)
```

Out[6]:
```
<torch._C.Generator at 0x158ef4b67f0>
```

In [7]:
```
# seed all available CUDA devices if using GPU
if torch.cuda.is_available():
    torch.cuda.manual_seed_all(seed)
```

Pipeline Configuration

In [8]:
```
# set pipeline mode to process training and validation images or test images, i.e., "train_val" for testing / validation and "test"
# for held-out test set
pipeline_mode = "train_val"
```

In [9]:
```
# set model name - BiomedCLIP model, available from Hugging Face through OpenCLIP
biomedclip_model_name = "hf-hub:microsoft/BiomedCLIP-PubMedBERT_256-vit_base_patch16_224"
```

In [10]:
```
# BiomedCLIP image embedding dimension
biomedclip_feature_dimension = 512
```

In [11]:
```
# determine whether pretrained BiomedCLIP visual model can be fine-tuned (for CUI classification)
freeze_pretrained_model = False
```

In [12]:
```
# include UMLS CUI concepts that appear in at least one ROCOv2 training imagege
minimum_concept_frequency = 1
```

In [13]:
```
# apply 20% dropout during training to help reduce overfitting
dropout = 0.2
```

In [14]:
```
# set GPU and CPU batch sizes (dependent on run-time environment)
batch_size = 16 if device.type == "cuda" else 4
```

In [15]:
```
# set number of training epochs for the concept classifier
num_epochs = 3
```

In [16]:
```
# set the learning rate used to update concept classifier weights during training
learning_rate = 0.00001
```

In [17]:
```
# apply a small penalty to large weights to help reduce overfitting
weight_decay = 0.0001
```

In [18]:
```
# stop training if validation loss does not improve for 2 consecutive epochs
patience = 2
```

In [19]:
```
# set number of PyTorch background worker processes used to load dataset batches
num_workers = 1
```

In [20]:
```
# Limit gradient magnitude during training to help prevent unstable weight updates
clip_grad_norm = 1.0
```

In [21]:
```
# retain predicted CUI concepts with a probability of at least 50%
minimum_prediction_probability = 0.5
```

In [22]:
```
# pipeline retention threshold, i.e. for output files
pipeline_retention_threshold = 0.8
```

In [23]:
```
# reuse the saved classifier checkpoint unless retraining is explicitly requested
force_retrain_concept_classifier = False #True
```

In [24]:
```
# set the conformal error level to 5%, i.e. the an acceptance threshold for each CUI concept
# to reduce false detections of that specific concept
conformal_alpha = 0.05
```

In [25]:
```
# require at least 20 calibration images where ground-truths do not
# contain a specific CUI. This ensures that there are enough negative examples to estimate
# the probability threshold used to accept or reject predictions for that CUI
minimum_calibration_examples = 20
```

In [26]:
```
# require symbolic rules derived from UMLS relationships
require_symbolic_rules = True
```

In [27]:
```
# set file and folder paths
training_images_folder = "ROCOv2/train_images/train"
validation_images_folder = "ROCOv2/valid_images/validation"
calibration_images_folder = "ROCOv2/calibration_images/calibration"
test_images_folder = "ROCOv2/test_images/test"
training_captions_filename = "ROCOv2/train_captions.csv"
validation_captions_filename = "ROCOv2/valid_captions_split.csv"
calibration_captions_filename = "ROCOv2/calibration_captions_split.csv"
test_captions_filename = "ROCOv2/test_captions.csv"
training_concepts_filename = "ROCOv2/train_concepts.csv"
validation_concepts_filename = "ROCOv2/valid_concepts.csv"
cui_mapping_filename = "ROCOv2/cui_mapping.csv"
concept_index_filename = "ROCOv2/biomedclip_concept_index.pt"
concept_classifier_checkpoint = "ROCOv2/biomedclip_concept_classifier_best.pth"
calibration_rules_filename = "ROCOv2/conformal_calibration_rules.csv"
pipeline_output_folder = "ROCOv2"
symbolic_rules_filename = "ROCOv2/rocov2_symbolic_rules_train_only.csv"
```

Load ROCOv2 Metadata

In [28]:
```
def load_required_csv(filename):
    """
    Load required CSV file

    Args:
        filename (str): CSV file path
```

```
    Returns:
        pd.DataFrame: loaded CSV data
    """

    try:

        # check file exists
        if not Path(filename).exists():
            raise FileNotFoundError(f"File not found: {filename}")

        # load CSV file
        return pd.read_csv(filename)

    except Exception as e:
        raise RuntimeError(f"Exception : {str(e)}")
```

In [29]:
```
# load CUI and concept name mapping files
cui_mapping = load_required_csv(cui_mapping_filename)
```

In [30]:
```
# load metadata required for selected pipeline (mode)
if pipeline_mode == "train_val":

    # load training and revised validation/calibration metadata
    training_captions = load_required_csv(training_captions_filename)
    validation_captions = load_required_csv(validation_captions_filename)
    calibration_captions = load_required_csv(calibration_captions_filename)
    training_concepts = load_required_csv(training_concepts_filename)
    original_validation_concepts = load_required_csv(validation_concepts_filename)

elif pipeline_mode == "test":

    # load test metadata used only to identify test images
    test_captions = load_required_csv(test_captions_filename)
```

Prepare Pipeline Datasets

In [31]:
```
# prepare image IDs and concept data for training, validation and calibration
if pipeline_mode == "train_val":

    # image IDs
```

```
    training_ids = training_captions["ID"].astype(str).drop_duplicates().tolist()
    validation_ids = validation_captions["ID"].astype(str).drop_duplicates().tolist()
    calibration_ids = calibration_captions["ID"].astype(str).drop_duplicates().tolist()

    # convert validation image IDs to strings to match against validation and calibration image IDs
    original_validation_concepts["ID"] = original_validation_concepts["ID"].astype(str)

    # retain concept labels for the validation set
    validation_concepts =
original_validation_concepts[original_validation_concepts["ID"].isin(validation_ids)].copy().reset_index(drop=True)

    # retain concept labels for the calibration set
    calibration_concepts =
original_validation_concepts[original_validation_concepts["ID"].isin(calibration_ids)].copy().reset_index(drop=True)

    # display split sizes
    print(f"Training images: {len(training_ids)}")
    print(f"Validation images: {len(validation_ids)}")
    print(f"Calibration images: {len(calibration_ids)}")
```

```
Training images: 59958
Validation images: 4952
Calibration images: 4952
```

In [32]:
```
# prepare test image IDs
if pipeline_mode == "test":

    # test image IDs
    test_ids = test_captions["ID"].astype(str).drop_duplicates().tolist()

    # display test split size
    print(f"Test images: {len(test_ids)}")
```

Stage 1 - Conceptual Pipeline

Concept Index

In [33]:
```
# create / load each UMLS CUI concept index mapping, i.e., to a classifier class number
if pipeline_mode == "train_val":

    # instantiate concept index
```

```
    concept_index = ConceptIndex()

    # build index from training CUIs
    concept_index.build_index(training_concepts, minimum_concept_frequency=minimum_concept_frequency)

    # add canonical UMLS concept names
    concept_index.add_concept_names(cui_mapping)

    # save concept index for test mode
    concept_index.save(concept_index_filename)

    # display number of CUI classes
    print(f"Concept classes: {len(concept_index)}")
```

```
Concept classes: 1947
```

In [34]:
```
if pipeline_mode == "test":

    # load concept index created during training and validation mode
    concept_index = ConceptIndex.load(concept_index_filename)

    # display number of CUI classes
    print(f"Concept classes: {len(concept_index)}")
```

BiomedCLIP ViT-B/16

In [35]:
```
# load pretrained BiomedCLIP model and image transformations
biomedclip_model, biomedclip_training_transform, biomedclip_validation_transform =
open_clip.create_model_and_transforms(biomedclip_model_name)
```

```
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits a
nd faster downloads.
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_
TOKEN to enable higher rate limits and faster downloads.
```

In [36]:
```
# only retain BiomedCLIP ViT-B/16 visual model
biomedclip_visual_model = biomedclip_model.visual
```

In [37]:
```
# remove unused text model reference
del biomedclip_model
```

In [38]:
```
# move pretrained visual model to selected device
biomedclip_visual_model = biomedclip_visual_model.to(device)
```

Concept Classification Datasets

In [39]:
```
# create data loaders for selected pipeline mode
if pipeline_mode == "train_val":

    # training dataset uses BiomedCLIP training transformation
    training_dataset = ROCOv2_ConceptDataset(training_images_folder, ".jpg", biomedclip_training_transform,
training_ids, concept_index,
                                             training_concepts)

    # deterministic training prediction dataset uses BiomedCLIP validation transformation
    training_prediction_dataset = ROCOv2_ConceptDataset(training_images_folder, ".jpg",
biomedclip_validation_transform, training_ids,
                                                        concept_index, training_concepts)

    # validation and calibration datasets use BiomedCLIP validation transformation
    validation_dataset = ROCOv2_ConceptDataset(validation_images_folder, ".jpg", biomedclip_validation_transform,
validation_ids,
                                               concept_index, validation_concepts)
    calibration_dataset = ROCOv2_ConceptDataset(calibration_images_folder, ".jpg", biomedclip_validation_transform,
calibration_ids,
                                                concept_index, calibration_concepts)

    # create data loaders
    training_loader = DataLoader(training_dataset, batch_size=batch_size, shuffle=True, num_workers=num_workers,
                                 pin_memory=(device.type == "cuda"))
    training_prediction_loader = DataLoader(training_prediction_dataset, batch_size=batch_size, shuffle=False,
                                            num_workers=num_workers, pin_memory=(device.type == "cuda"))
    validation_loader = DataLoader(validation_dataset, batch_size=batch_size, shuffle=False,
                                   num_workers=num_workers, pin_memory=(device.type == "cuda"))
    calibration_loader = DataLoader(calibration_dataset, batch_size=batch_size, shuffle=False,
num_workers=num_workers,
                                    pin_memory=(device.type == "cuda"))
```

In [40]:
```
if pipeline_mode == "test":

    # test dataset does not use test ground-truth concepts
```

```
    test_dataset = ROCOv2_ConceptDataset(test_images_folder, ".jpg", biomedclip_validation_transform, test_ids,
concept_index, concepts=None)

    # create test data loader
    test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers,
pin_memory=(device.type == "cuda"))
```

Concept Classifier

In [41]:
```
# instantiate BiomedCLIP multi-label CUI concept classifier
concept_classifier = ConceptClassifierModel(pretrained_model=biomedclip_visual_model,
feature_dimension=biomedclip_feature_dimension,
                                            num_concepts=len(concept_index), dropout=dropout,
freeze_pretrained_model=freeze_pretrained_model,
                                            feature_method="forward", pooling="mean").to(device)
```

In [42]:
```
# train or reload concept classifier when processing training and validation images
if pipeline_mode == "train_val":

    # reload saved best checkpoint
    if Path(concept_classifier_checkpoint).exists() and not force_retrain_concept_classifier:
        concept_classifier.load_checkpoint(concept_classifier_checkpoint, device)
        concept_classifier_history = concept_classifier.history
        print(f"Loaded concept-classifier checkpoint: {concept_classifier_checkpoint}")

    else:
        # calculate class weights to account for CUI class imbalance
        positive_weights = concept_index.calculate_positive_weights(training_concepts, training_ids)

        # train concept classifier and retain best validation checkpoint
        concept_classifier_history = concept_classifier.train_model(training_loader, validation_loader,
num_epochs=num_epochs,
                                                                     learning_rate=learning_rate,
weight_decay=weight_decay,
                                                                     patience=patience,
clip_grad_norm=clip_grad_norm,
                                                                     checkpoint_path=concept_classifier_checkpoint,
device=device,
                                                                     positive_weights=positive_weights)


        # confirm saved checkpoint
```

```
        print(f"Saved concept-classifier checkpoint: {concept_classifier_checkpoint}")

    # display training history
    display(concept_classifier_history)
```

```
Loaded concept-classifier checkpoint: ROCOv2/biomedclip_concept_classifier_best.pth
```

	epoch	training loss	validation loss	learning rate	overfitting
0	1	1.587408	1.896152	0.00001	0.308744
1	2	1.425542	1.799372	0.00001	0.373830

In [43]:
```
# load trained concept classifier when processing test images
if pipeline_mode == "test":

    # load best concept classifier checkpoint
    concept_classifier.load_checkpoint(concept_classifier_checkpoint, device)
```

Generate CUI Predictions

In [44]:
```
# generate CUI probability predictions
if pipeline_mode == "train_val":

    # generate training predictions
    training_predictions = concept_classifier.generate_predictions(training_prediction_loader, concept_index,
device,

minimum_probability=minimum_prediction_probability)

    # generate validation predictions
    validation_predictions = concept_classifier.generate_predictions(validation_loader, concept_index, device,

minimum_probability=minimum_prediction_probability)

    # generate calibration predictions for diagnostics
    calibration_predictions = concept_classifier.generate_predictions(calibration_loader, concept_index, device,

minimum_probability=minimum_prediction_probability)

    # display prediction counts
```

```
    print(f"Training prediction rows: {len(training_predictions)}")
    print(f"Validation prediction rows: {len(validation_predictions)}")
    print(f"Calibration prediction rows: {len(calibration_predictions)}")
```

```
Concept predictions:   0%|          | 0/14990 [00:08<?, ?it/s]
Concept predictions:   0%|          | 0/1238 [00:13<?, ?it/s]
Concept predictions:   0%|          | 0/1238 [00:07<?, ?it/s]
Training prediction rows: 10143590
Validation prediction rows: 904124
Calibration prediction rows: 903118
```

In [45]:

```python
if pipeline_mode == "train_val":

    # diagnose prediction counts at different probability thresholds
    diagnostic_thresholds = [0.50, 0.60, 0.70, 0.80, 0.90, 0.95]

    diagnostic_results = []

    for split_name, predictions, total_images in [
        ("Training", training_predictions, len(training_ids)),
        ("Validation", validation_predictions, len(validation_ids)),
        ("Calibration", calibration_predictions, len(calibration_ids))
    ]:

        for threshold in diagnostic_thresholds:

            # retain predictions meeting diagnostic threshold
            filtered_predictions = predictions[
                predictions["Probability CUI In The Image"] >= threshold
            ]

            # calculate predictions per image
            predictions_per_image = filtered_predictions.groupby("Image ID").size()

            # store diagnostic results
            diagnostic_results.append({
                "Split": split_name,
                "Threshold": threshold,
                "Rows": len(filtered_predictions),
                "Average CUIs Per Image": len(filtered_predictions) / total_images,
                "Median CUIs Per Image": predictions_per_image.median() if not predictions_per_image.empty else 0,
                "Maximum CUIs Per Image": predictions_per_image.max() if not predictions_per_image.empty else 0
```

```
        })

    # create diagnostic table
    prediction_diagnostics = pd.DataFrame(diagnostic_results)

    # display results
    display(prediction_diagnostics)
```

	Split	Threshold	Rows	Average CUIs Per Image	Median CUIs Per Image	Maximum CUIs Per Image
0	Training	0.50	10143590	169.178258	157.0	999
1	Training	0.60	6670666	111.255646	101.0	722
2	Training	0.70	4101059	68.398863	65.0	457
3	Training	0.80	2164441	36.099286	30.0	247
4	Training	0.90	663189	11.060893	9.0	79
5	Training	0.95	185366	3.091597	3.0	39
6	Validation	0.50	904124	182.577544	170.0	942
7	Validation	0.60	589678	119.078756	107.0	658
8	Validation	0.70	359517	72.600363	68.0	410
9	Validation	0.80	187583	37.880250	32.0	226
10	Validation	0.90	56686	11.447092	10.0	75
11	Validation	0.95	15363	3.102383	3.0	29
12	Calibration	0.50	903118	182.374394	167.0	895
13	Calibration	0.60	590902	119.325929	107.0	630
14	Calibration	0.70	361054	72.910743	69.0	392
15	Calibration	0.80	189523	38.272011	34.0	209
16	Calibration	0.90	58336	11.780291	10.0	66
17	Calibration	0.95	15658	3.161955	3.0	28

In [46]:
```
if pipeline_mode == "test":

    # generate test predictions
    test_predictions = concept_classifier.generate_predictions(test_loader, concept_index, device,
minimum_probability=minimum_prediction_probability)

    # display prediction count
    print(f"Test prediction rows: {len(test_predictions)}")
```

Stage 2 - Neuro-Symbolic Pipeline

UMLS Symbolic Knowledge

In [47]:
```
# instantiate UMLS symbolic validator
symbolic_validator = UMLSSymbolicValidator(concept_index)
```

In [48]:
```
# generate UMLS symbolic rules using training concepts
if pipeline_mode == "train_val":

    # initialise UMLS symbolic rule generator
    umls_rule_generator = UMLSCUISymbolicRuleGenerator(concept_index=concept_index)

    # generate training-only rules using direct and two-step UMLS relationships
    symbolic_rules = umls_rule_generator.create_automatic_training_rules(
        training_concepts=training_concepts,
        output_folder="ROCOv2",
        minimum_support=10,
        minimum_confidence=0.98,
        minimum_lift=1.5,
        require_umls_relationship=True,
        umls_mrrel_filename="ROCOv2/MRREL.RRF",
        maximum_umls_hops=2
    )
```

```
Training-derived candidate rules: 61
Candidates with direct UMLS links: 1
Candidates with hierarchical two-step UMLS links: 2
Generated symbolic rules: 3
Symbolic rules saved to: ROCOv2\rocov2_symbolic_rules_train_only.csv
Rule evidence saved to: ROCOv2\rocov2_symbolic_rule_evidence_train_only.csv
```

In [49]:
```
# load symbolic rules
symbolic_validator.load_rules(symbolic_rules_filename)
```

Out[49]:
```
<conceptual_neurosymbolic_conformal_pipeline.UMLSSymbolicValidator at 0x158858d96a0>
```

In [50]:
```
if pipeline_mode == "train_val":

    # validate training predictions
    training_symbolic_predictions = symbolic_validator.validate_predictions(training_predictions)

    # validate validation predictions
    validation_symbolic_predictions = symbolic_validator.validate_predictions(validation_predictions)

    # validate calibration predictions
    calibration_symbolic_predictions = symbolic_validator.validate_predictions(calibration_predictions)
```

In [51]:
```
if pipeline_mode == "test":

    # validate test predictions
    test_symbolic_predictions = symbolic_validator.validate_predictions(test_predictions)
```

Stage 3 - Conformal Pipeline

Conformal Calibration

In [52]:
```
# create or load conformal calibration rules
if pipeline_mode == "train_val":

    # instantiate conformal calibrator
    conformal_calibrator = ConformalCalibrator(alpha=conformal_alpha,
minimum_calibration_examples=minimum_calibration_examples)

    # create CUI-specific conformal thresholds from held-out calibration images
    calibration_rules = conformal_calibrator.fit(concept_classifier, calibration_loader, concept_index, device)

    # save calibration rules for test mode
    conformal_calibrator.save(calibration_rules_filename)
```

```
    # display calibration rules
    display(calibration_rules.head(10))
```

```
Conformal calibration:   0%|          | 0/1238 [00:12<?, ?it/s]
```

	CUI	Conformal Threshold	Calibration Examples	Alpha	Calibrated
0	C0000726	0.828449	4656	0.05	True
1	C0000833	0.657009	4938	0.05	True
2	C0000846	0.777609	4949	0.05	True
3	C0000962	0.886648	4947	0.05	True
4	C0001074	0.915843	4949	0.05	True
5	C0001162	0.448736	4952	0.05	True
6	C0001168	0.937818	4925	0.05	True
7	C0001208	0.790555	4950	0.05	True
8	C0001209	0.673889	4950	0.05	True
9	C0001304	0.728966	4929	0.05	True

In [53]:
```
if pipeline_mode == "test":

    # load calibration rules created using held-out calibration images
    conformal_calibrator = ConformalCalibrator.load(calibration_rules_filename)
```

Apply Conformal Rules

In [54]:
```
# apply conformal rules
if pipeline_mode == "train_val":

    # apply rules to symbolically validated training predictions
    training_conformal_predictions = conformal_calibrator.apply(training_symbolic_predictions)

    # apply rules to symbolically validated validation predictions
    validation_conformal_predictions = conformal_calibrator.apply(validation_symbolic_predictions)
```

In [55]:
```
if pipeline_mode == "test":

    # apply saved rules to symbolically validated test predictions
    test_conformal_predictions = conformal_calibrator.apply(test_symbolic_predictions)
```

Combined Concept Reliability

In [56]:
```
def add_reliability(predictions):
    """
    Assign final concept reliability

    Args:
        predictions (DataFrame): predictions containing symbolic and conformal validity

    Returns:
        pd.DataFrame: predictions with Reliable flag
    """

    try:

        # copy predictions
        reliable_predictions = predictions.copy()

        # mark a CUI prediction as reliable only if it passes both the symbolic rules
        # and the conformal probability threshold
        reliable_predictions["Reliable"] = reliable_predictions["Symbolic (Valid)"] &
reliable_predictions["Conformal (Valid)"]

        return reliable_predictions

    except Exception as e:
        raise RuntimeError(f"Exception : {str(e)}")
```

In [57]:
```
# assign final reliability rating
if pipeline_mode == "train_val":

    # training reliability
    training_reliable_predictions = add_reliability(training_conformal_predictions)
```

```
    # validation reliability
    validation_reliable_predictions = add_reliability(validation_conformal_predictions)
```

In [58]:
```
if pipeline_mode == "test":

    # test reliability
    test_reliable_predictions = add_reliability(test_conformal_predictions)
```

Create Pipeline DataFrames

In [59]:
```
# instantiate pipeline controller
pipeline_controller = ConceptualNeuroSymbolicConformalPipeline(mode=pipeline_mode)
```

In [60]:
```
# create final pipeline DataFrames
if pipeline_mode == "train_val":

    # create training pipeline
    training_pipeline = pipeline_controller.create_pipeline(training_predictions, symbolic_validator,
conformal_calibrator)

    # create validation pipeline
    validation_pipeline = pipeline_controller.create_pipeline(validation_predictions, symbolic_validator,
conformal_calibrator)
```

In [61]:
```
if pipeline_mode == "test":

    # create test pipeline
    test_pipeline = pipeline_controller.create_pipeline(test_predictions, symbolic_validator, conformal_calibrator)
```

Export Pipeline Files

In [62]:
```
# export training and validation pipeline CSV files
if pipeline_mode == "train_val":

    # export complete training pipeline including rejected predictions
    training_pipeline_filename = pipeline_controller.save_pipeline(
        training_pipeline,
        pipeline_output_folder,
        "training"
```

```
)

# export complete validation pipeline including rejected predictions
validation_pipeline_filename = pipeline_controller.save_pipeline(
    validation_pipeline,
    pipeline_output_folder,
    "validation"
)

# select training predictions meeting the retention threshold
training_retained_predictions = training_predictions[
    training_predictions["Probability CUI In The Image"] >= pipeline_retention_threshold
].copy()

# reapply symbolic and conformal validation to retained training predictions
training_retained_pipeline = pipeline_controller.create_pipeline(
    training_retained_predictions,
    symbolic_validator,
    conformal_calibrator
)

# retain reliable training predictions
training_accepted_pipeline = training_retained_pipeline[
    training_retained_pipeline["Reliable"]
].reset_index(drop=True)

# save accepted training predictions separately
training_accepted_filename = Path(pipeline_output_folder) / "pipeline_training_accepted.csv"
training_accepted_pipeline.to_csv(training_accepted_filename, index=False)

# select validation predictions meeting the retention threshold
validation_retained_predictions = validation_predictions[
    validation_predictions["Probability CUI In The Image"] >= pipeline_retention_threshold
].copy()

# reapply symbolic and conformal validation to retained validation predictions
validation_retained_pipeline = pipeline_controller.create_pipeline(
    validation_retained_predictions,
    symbolic_validator,
    conformal_calibrator
)
```

```
    # retain reliable validation predictions
    validation_accepted_pipeline = validation_retained_pipeline[
        validation_retained_pipeline["Reliable"]
    ].reset_index(drop=True)

    # save accepted validation predictions separately
    validation_accepted_filename = Path(pipeline_output_folder) / "pipeline_validation_accepted.csv"
    validation_accepted_pipeline.to_csv(validation_accepted_filename, index=False)

    # display saved files
    print(training_pipeline_filename)
    print(validation_pipeline_filename)
    print(training_accepted_filename)
    print(validation_accepted_filename)
```

```
ROCOv2\pipeline_training.csv
ROCOv2\pipeline_validation.csv
ROCOv2\pipeline_training_accepted.csv
ROCOv2\pipeline_validation_accepted.csv
```

In [63]:

```
# export test pipeline CSV files
if pipeline_mode == "test":

    # export complete test pipeline including rejected predictions
    test_pipeline_filename = pipeline_controller.save_pipeline(
        test_pipeline,
        pipeline_output_folder,
        "test"
    )

    # select test predictions meeting the retention threshold
    test_retained_predictions = test_predictions[
        test_predictions["Probability CUI In The Image"] >= pipeline_retention_threshold
    ].copy()

    # reapply symbolic and conformal validation to retained test predictions
    test_retained_pipeline = pipeline_controller.create_pipeline(
        test_retained_predictions,
        symbolic_validator,
        conformal_calibrator
    )
```

```
    # retain reliable test predictions
    test_accepted_pipeline = test_retained_pipeline[
        test_retained_pipeline["Reliable"]
    ].reset_index(drop=True)

    # save accepted test predictions separately
    test_accepted_filename = Path(pipeline_output_folder) / "pipeline_test_accepted.csv"
    test_accepted_pipeline.to_csv(test_accepted_filename, index=False)

    # display saved files
    print(test_pipeline_filename)
    print(test_accepted_filename)
```

Validate Pipeline Outputs

In [64]:

```
def validate_pipeline_output(pipeline, expected_image_ids, split_name):
    """
    Validate final pipeline structure and values

    Args:
        pipeline (DataFrame): final pipeline data
        expected_image_ids: image IDs belonging to the split
        split_name (str): display name

    Returns:
        dict: validation summary
    """

    try:

        # required output columns
        required_columns = ["Image ID", "CUI", "Concept Name", "Probability CUI In The Image", "Symbolic (Valid)",
"Conformal (Valid)", "Reliable"]

        # check required columns
        if list(pipeline.columns) != required_columns:
            raise ValueError(f"{split_name} pipeline columns do not match the required output structure")

        # check probabilities
        if not pipeline["Probability CUI In The Image"].between(0, 1).all():
            raise ValueError(f"{split_name} pipeline contains probabilities outside 0 to 1")
```

```
        # check duplicate image/CUI rows
        if pipeline.duplicated(subset=["Image ID", "CUI"]).any():
            raise ValueError(f"{split_name} pipeline contains duplicate image/CUI rows")

        # check image IDs belong to expected split
        unexpected_image_ids = set(pipeline["Image ID"].astype(str)) - set(str(image_id) for image_id in
expected_image_ids)
        if unexpected_image_ids:
            raise ValueError(f"{split_name} pipeline contains unexpected image IDs")

        # check reliability logic
        expected_reliable = pipeline["Symbolic (Valid)"] & pipeline["Conformal (Valid)"]
        if not pipeline["Reliable"].equals(expected_reliable):
            raise ValueError(f"{split_name} pipeline contains invalid Reliable values")

        # create validation summary
        validation_summary = {"Split": split_name, "Expected Images": len(set(expected_image_ids)), "Images With
Candidate Concepts": pipeline["Image ID"].nunique(), "Rows": len(pipeline), "Unique CUIs":
pipeline["CUI"].nunique()}

        return validation_summary

    except Exception as e:
        raise RuntimeError(f"Exception : {str(e)}")
```

In [65]:
```
# validate pipeline outputs
validation_results = []
```

In [66]:
```
if pipeline_mode == "train_val":

    # validate training and validation pipeline outputs
    validation_results.append(validate_pipeline_output(training_pipeline, training_ids, "Training"))
    validation_results.append(validate_pipeline_output(validation_pipeline, validation_ids, "Validation"))
```

In [67]:
```
if pipeline_mode == "test":

    # validate test pipeline output
    validation_results.append(validate_pipeline_output(test_pipeline, test_ids, "Test"))
```

In [68]:
```
# display validation results
display(pd.DataFrame(validation_results))
```

	Split	Expected Images	Images With Candidate Concepts	Rows	Unique CUIs
0	Training	59958	59958	10143590	1809
1	Validation	4952	4952	904124	1755

Pipeline Statistics

In [69]:
```
def calculate_pipeline_statistics(pipeline, split_name):
    """
    Calculate pipeline acceptance statistics

    Args:
        pipeline (DataFrame): final pipeline data
        split_name (str): display name

    Returns:
        dict: pipeline statistics
    """

    try:

        # calculate counts
        prediction_rows = len(pipeline)
        symbolic_valid = int(pipeline["Symbolic (Valid)"].sum()) if prediction_rows > 0 else 0
        conformal_valid = int(pipeline["Conformal (Valid)"].sum()) if prediction_rows > 0 else 0
        reliable = int(pipeline["Reliable"].sum()) if prediction_rows > 0 else 0

        # calculate percentages
        symbolic_acceptance = (symbolic_valid / prediction_rows) * 100 if prediction_rows > 0 else 0.0
        conformal_acceptance = (conformal_valid / prediction_rows) * 100 if prediction_rows > 0 else 0.0
        reliable_percentage = (reliable / prediction_rows) * 100 if prediction_rows > 0 else 0.0

        # return statistics
        return {"Split": split_name, "Prediction Rows": prediction_rows, "Images": pipeline["Image ID"].nunique(),
"Unique CUIs": pipeline["CUI"].nunique(), "Mean Probability": pipeline["Probability CUI In The Image"].mean() if
prediction_rows > 0 else 0.0, "Symbolic Valid": symbolic_valid, "Symbolic Acceptance %": symbolic_acceptance,
```

```
"Conformal Valid": conformal_valid, "Conformal Acceptance %": conformal_acceptance, "Reliable": reliable, "Reliable %": reliable_percentage}

    except Exception as e:
        raise RuntimeError(f"Exception : {str(e)}")
```

In [70]:
```
# calculate pipeline statistics
pipeline_statistics = []
```

In [71]:
```
if pipeline_mode == "train_val":

    # calculate training and validation statistics
    pipeline_statistics.append(calculate_pipeline_statistics(training_pipeline, "Training"))
    pipeline_statistics.append(calculate_pipeline_statistics(validation_pipeline, "Validation"))
```

In [72]:
```
if pipeline_mode == "test":

    # calculate test statistics
    pipeline_statistics.append(calculate_pipeline_statistics(test_pipeline, "Test"))
```

In [73]:
```
# display pipeline statistics
display(pd.DataFrame(pipeline_statistics))
```

	Split	Prediction Rows	Images	Unique CUIs	Mean Probability	Symbolic Valid	Symbolic Acceptance %	Conformal Valid	Conformal Acceptance %	Reliable	Reliable %
0	Training	10143590	59958	1809	0.680786	10142710	99.991325	3631301	35.798973	3631194	35.797918
1	Validation	904124	4952	1755	0.678625	904021	99.988608	330083	36.508598	330072	36.507382

References

BiomedCLIP-PubMedBERT_256-vit_base_patch16_224: https://huggingface.co/microsoft/BiomedCLIP-PubMedBERT_256-vit_base_patch16_224

OCOM5100M Programming For Data Science

OCOM5101M Data Science

OCOM5200M Machine Learning

OCOM5203M Deep Learning

OCOM5203M Deep Learning Module, Summative Assessment 2

OCOM5204M Data Mining And Analytics

Project Jupyter, Jupyter Notebook documentation. Available At: https://jupyter-notebook.readthedocs.io/en/stable/

Python Software Foundation, Python 3.13 documentation (2026). Available At: https://docs.python.org/3.13/

ROCOv2: Radiology Objects in Context Version 2, An Updated Multimodal Image Dataset: https://doi.org/10.1038/s41597-024-03496-6